References

AI Fraud Detector for Small Businesses · generated November 09, 2026

- Friedman, J. H. (2001). Greedy function approximation: a gradient boosting machine. *Annals of Statistics*, 29(5).
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions (SHAP). *NeurIPS*.
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why should I trust you?" Explaining the predictions of any classifier (LIME). *KDD*.
- Dal Pozzolo, A., et al. (2015). Calibrating probability with undersampling for unbalanced classification. *IEEE SSCI*.
- Saito, T., & Rehmsmeier, M. (2015). The precision-recall plot is more informative than the ROC plot on imbalanced datasets. *PLOS ONE*.
- Yurdakul, B. (2018). Statistical properties of population stability index. Western Michigan University.
- Pedregosa, F., et al. (2011). Scikit-learn: machine learning in Python. *JMLR* 12.
- PCI Security Standards Council. PCI DSS v4.0 - guidance on tokenisation and storage of cardholder data.

Note: the dataset in 2_Data is synthetic, generated by 2_Data/generate_data.py to mimic published fraud patterns (fraud rings, card testing, account takeover). Replace it with your store's labelled history before production use.